

flights_example_DataFrame

October 5, 2020

```
[3]: from pyspark import sql
      from pyspark.sql import functions as f, udf

      sqlContext = sql.SparkSession.builder \
        .master("local") \
        .appName("Flight DF") \
        .getOrCreate()

      flights = sqlContext.read.format('csv')\
        .options(header='true', inferSchema='true')\
        .load("flights.csv.bz2")

      airport = sqlContext.read.format('csv')\
        .options(header='true', inferSchema='true')\
        .load("airports.csv.bz2")
```

```
[43]: linesHeader = lines.first()
      flights_raw = lines\
        .zipWithIndex()\
        .filter(lambda x: x[1] > 2)\
        .keys()

      flights = flights_raw\
        .map(lambda x: x.split(','))\
        .map(lambda x: (x[0], x[1], x[2], x[3], x[4],
                        x[5], x[6], x[7], x[8], x[9],
                        x[10], 0 if x[11]==' ' else float(x[11]), x[12], x[13],
↪x[14],
                        x[15], x[16], x[17], x[18], x[19],
                        x[20], x[21], x[22], x[23], x[24],
                        x[25], x[26], x[27], x[28], x[29], x[30])
      ))

      airports_raw = airports_lines\
        .zipWithIndex()\
```

```

        .filter(lambda x:x[1]>2)\
        .keys()
airports = airports_raw\
    .map(lambda x: x.split(','))
mainFlightsData = flights.map(lambda p:
                                (p[0], p[1], p[2], p[3], p[4], p[5], p[6],
                                 p[7], p[8], p[9], p[10], p[11], p[24]))

```

```

↳ -----
NameError                                Traceback (most recent call↳
↳last)

```

```

<ipython-input-43-b883e54e11f3> in <module>
----> 1 linesHeader = lines.first()
      2 flights_raw = lines\
      3     .zipWithIndex()\
      4     .filter(lambda x: x[1] > 2)\
      5     .keys()

```

NameError: name 'lines' is not defined

```

[4]: flights.printSchema()
     airport.printSchema()

```

```

root
|-- YEAR: integer (nullable = true)
|-- MONTH: integer (nullable = true)
|-- DAY: integer (nullable = true)
|-- DAY_OF_WEEK: integer (nullable = true)
|-- AIRLINE: string (nullable = true)
|-- FLIGHT_NUMBER: integer (nullable = true)
|-- TAIL_NUMBER: string (nullable = true)
|-- ORIGIN_AIRPORT: string (nullable = true)
|-- DESTINATION_AIRPORT: string (nullable = true)
|-- SCHEDULED_DEPARTURE: integer (nullable = true)
|-- DEPARTURE_TIME: integer (nullable = true)
|-- DEPARTURE_DELAY: integer (nullable = true)
|-- TAXI_OUT: integer (nullable = true)
|-- WHEELS_OFF: integer (nullable = true)
|-- SCHEDULED_TIME: integer (nullable = true)
|-- ELAPSED_TIME: integer (nullable = true)
|-- AIR_TIME: integer (nullable = true)

```

```

|-- DISTANCE: integer (nullable = true)
|-- WHEELS_ON: integer (nullable = true)
|-- TAXI_IN: integer (nullable = true)
|-- SCHEDULED_ARRIVAL: integer (nullable = true)
|-- ARRIVAL_TIME: integer (nullable = true)
|-- ARRIVAL_DELAY: integer (nullable = true)
|-- DIVERTED: integer (nullable = true)
|-- CANCELLED: integer (nullable = true)
|-- CANCELLATION_REASON: string (nullable = true)
|-- AIR_SYSTEM_DELAY: integer (nullable = true)
|-- SECURITY_DELAY: integer (nullable = true)
|-- AIRLINE_DELAY: integer (nullable = true)
|-- LATE_AIRCRAFT_DELAY: integer (nullable = true)
|-- WEATHER_DELAY: integer (nullable = true)

```

root

```

|-- IATA_CODE: string (nullable = true)
|-- AIRPORT: string (nullable = true)
|-- CITY: string (nullable = true)
|-- STATE: string (nullable = true)
|-- COUNTRY: string (nullable = true)
|-- LATITUDE: double (nullable = true)
|-- LONGITUDE: double (nullable = true)

```

```

[7]: #Q1 Find a list of Origin Airports
      flights.select("ORIGIN_AIRPORT").distinct().show()

```

```

+-----+
|ORIGIN_AIRPORT|
+-----+
|          BGM|
|          PSE|
|          INL|
|          DLG|
|        12888|
|          MSY|
|          PPG|
|        12003|
|        15041|
|          GEG|
|          SNA|
|          BUR|
|          GRB|
|          GTF|
|        14986|
|        13851|
|          IDA|

```

```
|          11150|
|          15412|
|           GRR|
+-----+
only showing top 20 rows
```

```
[8]: #Q2 Find a list of (Origin, Destination) pairs
      flights.select("ORIGIN_AIRPORT", "DESTINATION_AIRPORT").distinct().show()
```

```
+-----+-----+
|ORIGIN_AIRPORT|DESTINATION_AIRPORT|
+-----+-----+
|          BQN|          MCO|
|          PHL|          MCO|
|          MCI|          IAH|
|          SPI|          ORD|
|          SNA|          PHX|
|          LBB|          DEN|
|          ORD|          PDX|
|          EWR|          STT|
|          ATL|          GSP|
|          MCI|          MKE|
|          PBI|          DCA|
|          SMF|          BUR|
|          MDW|          MEM|
|          LAS|          LIT|
|          TPA|          ACY|
|          DSM|          EWR|
|          FSD|          ATL|
|          SJC|          LIH|
|          CLE|          SJU|
|        11298|        11057|
+-----+-----+
only showing top 20 rows
```

```
[10]: #Q3 Find the Origin airport which had the largest departure delay in the month
      ↪ of January
      flights.where(flights.MONTH == 1)\
        .orderBy("DEPARTURE_DELAY", ascending=False)\
        .limit(1)\
        .select("ORIGIN_AIRPORT")\
        .show()
```

```
+-----+
|ORIGIN_AIRPORT|
+-----+
```

	BHM

[11]: #Q4 Find out which carrier has the largest delay on Weekends.

```
flights.filter("DAY_OF_WEEK = 6 OR DAY_OF_WEEK = 7")\
  .orderBy("DEPARTURE_DELAY", ascending=False)\
  .limit(1)\
  .select("AIRLINE")\
  .show()
```

AIRLINE
AA

[12]: #Q5 Which airport has the most cancellation of flights?

```
flights.filter("CANCELLED = 1")\
  .withColumn("COUNT", f.lit(1))\
  .groupBy("ORIGIN_AIRPORT")\
  .agg(f.sum("COUNT").alias("COUNT"))\
  .orderBy("COUNT", ascending=False)\
  .limit(1)\
  .select("ORIGIN_AIRPORT", "COUNT")\
  .show()
```

ORIGIN_AIRPORT	COUNT
ORD	8548

[14]: #Q6 Find the percent of flights cancelled for each carrier.

```
flights.withColumn("TOTAL", f.lit(1))\
  .groupBy("AIRLINE")\
  .agg(f.sum("CANCELLED").alias("TOTAL_CANCELLED"), f.sum("TOTAL").\
    alias("TOTAL"))\
  .withColumn("CANCEL_RATE", f.col("TOTAL_CANCELLED")/f.col("TOTAL")*100)\
  .show()
```

AIRLINE	TOTAL_CANCELLED	TOTAL	CANCEL_RATE
UA	6573	515723	1.274521400053905

NK	2004	117379	1.7072900604026275
AA	10919	725984	1.5040276369727157
EV	15231	571977	2.6628693111785964
B6	4276	267048	1.6012102693148795
DL	3824	875881	0.4365889886868193
OO	9960	588353	1.6928612584621818
F9	588	90836	0.6473204456382932
US	4067	198715	2.0466497244797823
MQ	15025	294632	5.0995818512585185
HA	171	76272	0.22419760855884205
AS	669	172521	0.38777887909298
VX	534	61903	0.8626399366751207
WN	16043	1261855	1.2713822111098343

[16]: #Q7 Find the largest departure delay for each carrier

```
flights.groupBy("AIRLINE")\
  .agg(f.max("DEPARTURE_DELAY").alias("MAX_DEPARTURE_DELAY"))\
  .show()
```

AIRLINE	MAX_DEPARTURE_DELAY
UA	1314
NK	836
AA	1988
EV	1274
B6	1006
DL	1289
OO	1378
F9	1112
US	759
MQ	1544
HA	1433
AS	963
VX	644
WN	665

[17]: #Q8 Find the largest departure delay for each carrier for each month

```
flights.groupBy("AIRLINE", "MONTH")\
  .agg(f.max("DEPARTURE_DELAY").alias("MAX_DEPARTURE_DELAY"))\
  .show()
```

AIRLINE	MONTH	MAX_DEPARTURE_DELAY
NK	11	476
VX	10	430
UA	12	1194
HA	10	1022
OO	3	874
OO	4	878
OO	9	893
F9	2	852
F9	12	781
HA	5	326
UA	4	1314
MQ	10	1544
HA	12	1095
EV	4	757
DL	6	1201
DL	3	1166
DL	8	1207
B6	6	507
DL	10	1120
OO	10	1122

only showing top 20 rows

```
[19]: #Q9 For each carrier find the average Departure delay
flights.withColumn("TOTAL", f.lit(1))\
    .groupBy("AIRLINE")\
    .agg(f.sum("DEPARTURE_DELAY").alias("TOTAL_DEPARTURE_DELAY"), f.
    ↪sum("TOTAL").alias("TOTAL"))\
    .withColumn("AVG_DEPARTURE_DELAY", f.col("TOTAL_DEPARTURE_DELAY")/f.
    ↪col("TOTAL"))\
    .show()
```

AIRLINE	TOTAL_DEPARTURE_DELAY	TOTAL	AVG_DEPARTURE_DELAY
UA	7355348	515723	14.26220664969373
NK	1840887	117379	15.68327383944317
AA	6369435	725984	8.773519802089302
EV	4857338	571977	8.49219111957299
B6	3026467	267048	11.333044995656211
DL	6427294	875881	7.338090448359994
OO	4517510	588353	7.67823058605973
F9	1205449	90836	13.27060856928971
US	1196447	198715	6.0209194071912036
MQ	2837908	294632	9.63204268375465

HA	36972	76272	0.48473882945248586
AS	306997	172521	1.7794761217474975
VX	553852	61903	8.947094648078446
WN	13186520	1261855	10.450107183471951

```
[20]: #Q10 For each carrier find the average Departure delay for each month
flights.withColumn("TOTAL", f.lit(1))\
    .groupBy("AIRLINE", "MONTH")\
    .agg(f.sum("DEPARTURE_DELAY").alias("TOTAL_DEPARTURE_DELAY"), f.
    ↪sum("TOTAL").alias("TOTAL"))\
    .withColumn("AVG_DEPARTURE_DELAY", f.col("TOTAL_DEPARTURE_DELAY")/f.
    ↪col("TOTAL"))\
    .select("AIRLINE", "MONTH", "AVG_DEPARTURE_DELAY")
    .show()
```

AIRLINE	MONTH	TOTAL_DEPARTURE_DELAY	TOTAL	AVG_DEPARTURE_DELAY
NK	11	87001	10164	8.559720582447856
VX	10	38540	5464	7.053440702781844
UA	12	761043	43443	17.51819625716456
HA	10	1049	6242	0.1680551105414931
OO	3	289928	50078	5.789528335796158
OO	4	260302	49329	5.276855399460763
OO	9	182835	47625	3.8390551181102364
F9	2	146727	5809	25.258564296780857
F9	12	129059	8120	15.89396551724138
HA	5	-8676	6434	-1.3484612993472178
UA	4	532506	41342	12.880508925547868
MQ	10	75123	21982	3.4174779364934946
HA	12	-2771	6260	-0.4426517571884984
EV	4	328999	49296	6.673949204803635
DL	6	837824	77255	10.844916186654585
DL	3	622004	74166	8.386646172100424
DL	8	626586	80947	7.74069452851866
B6	6	255272	22558	11.31625144073056
DL	10	242914	75552	3.2151895383312157
OO	10	177362	48808	3.6338714964759875

only showing top 20 rows

```
[22]: #Q11 Which date of year has the highest rate of flight cancellations?
# Rate of flight cancellation is calculated by deviding number of canceled
    ↪flights by total number of flights.
```



```

flights.withColumn("TOTAL", f.lit(1))\
    .groupBy("YEAR", "MONTH", "DAY")\
    .agg(f.sum("CANCELLED").alias("TOTAL_CANCELLED"), f.sum("TOTAL").
↪alias("TOTAL"))\
    .withColumn("CANCEL_RATE", f.col("TOTAL_CANCELLED")/f.col("TOTAL")*100)\
    .orderBy("CANCEL_RATE", ascending=False)\
    .limit(1)\
    .select("YEAR", "MONTH", "DAY")\
    .show()

```

```

+----+-----+----+
|YEAR|MONTH|DAY|
+----+-----+----+
|2015|    1| 27|
+----+-----+----+

```

[48]: *#Q12 Calculate the number of flights to each destination state
For each carrier, for which state do they have the largest average delay?
You will need the airline and airport data sets for this question.*

```

from pyspark.sql.types import ArrayType, IntegerType, StringType
from pyspark.sql.functions import udf
fold_list = udf(lambda x,y: sorted(zip(x,y))[-1][1],StringType())

#Q8 Find the largest departure delay for each carrier for each month
flights.withColumn("COUNT", f.lit(1))\
    .groupBy("AIRLINE", "DESTINATION_AIRPORT")\
    .agg(f.sum("DEPARTURE_DELAY").alias("TOTAL_DEPARTURE_DELAY"),f.sum("COUNT").
↪alias("COUNT"))\
    .withColumn("DEPARTURE_AVG_DELAY", f.col("TOTAL_DEPARTURE_DELAY")/f.
↪col("COUNT"))\
    .join(airport, flights.DESTINATION_AIRPORT == airport.IATA_CODE)\
    .select("AIRLINE", "TOTAL_DEPARTURE_DELAY", "STATE")\
    .groupBy("AIRLINE")\
    .agg(
        f.collect_list("TOTAL_DEPARTURE_DELAY").alias("delay"),
        f.collect_list("STATE").alias("state")
    )\
    .withColumn("MAX_AVGDELAY_STATE", fold_list(f.col("delay"), f.
↪col("state")))\
    .select("AIRLINE", "MAX_AVGDELAY_STATE")\
    .show()

```

```

+-----+-----+
|AIRLINE|MAX_AVGDELAY_STATE|

```

	UA	IL
	NK	IL
	AA	TX
	EV	GA
	B6	NY
	DL	GA
	OO	IL
	F9	CO
	US	NC
	MQ	IL
	HA	HI
	AS	WA
	VX	CA
	WN	IL

[]: